

GPU并行计算和 CUDA程序开发及优化

周 斌 2014年秋

课程概况

▶ 时间

- ▶ 7~16周 周六 上午 9:45~12:10
- ▶ 实验课程 9~16周 周六下午 14:00 – 17:00; 晚18:00-21:00;

▶ 地点

- ▶ 课程:西 区3教副楼3C121
- ▶ 实验: 电四楼209

▶ 人物

- ▶ 老师+助教+同学

▶ 事件

- ▶ 课程+实验+....



计划增加一次课程

▶ 时间

- ▶ 周日 上午 9:50~11:20 及 下午 2:00~3:30

▶ 地点

- ▶ 课程:西区3教副楼3C121 或待定

▶ 人物

- ▶ 老师+助教+同学

▶ 事件

- ▶ 课程.



授课老师



▶ 周 斌

- ▶ 山东省科学院海仪所海洋遥感遥测研究室主任
- ▶ NVIDIA CUDA Fellow、中科大客座研究员
- ▶ 曾任NVIDIA 高性能计算开发技术资深工程师
- ▶ 科研:信号和信息处理、视频图像处理
- ▶ 兴趣:数据分析, 密码学和密码分析, 无人机系统
- ▶ 涉及领域:
 - ▶ 气象数值预报系统、视频图像处理、高速信号处理、生物信息学、网络搜索、加密系统等
- ▶ 关注GPU系统架构、高性能应用、移动互联网等

授课老师



▶ 邮件:

▶ synosy@gmail.com

▶ 答疑:

▶ 课上课下

▶ QQ: GPU深度开发

▶ 邮件

▶ [iqiyi 搜索 CUDA 教程](#)

▶ 风格: 中英夹杂



名人堂




李索恒



陈凯



曹旭磊



????

2012年秋季学期：李索恒

2013年春季学期：陈 凯

2013年秋季学期：曹旭磊

成绩优秀的同学获得了NVIDIA公司提供的奖励
本学期奖励继续么？



课程情况总结

▶ 2012年秋

- ▶ 90人选课，平均成绩83分，除1人退课外，全部及格并基本符合正态分布
- ▶ 12个Project
 - ▶ 内容: MD5&SHA-1; FDTD; Ray-tracing; Page-Rank; Sparse-Solver; Smith-Waterman....
- ▶ 3人获得NVIDIA公司的奖励

▶ 2013年春

- ▶ 23人选课，平均成绩85分，全部及格
 - ▶ 1人获得NVIDIA公司的奖励
-



课程情况总结

▶ 2013年秋

- ▶ 168人选课，平均成绩85分，12人退课，全部及格并基本符合正态分布；20个Project
- ▶ 3人获得NVIDIA公司的奖励



The Students

- ▶ Gain the abilities of:
 - ▶ Construction of GPU hardware systems
 - ▶ CUDA/GPU Programming for their research work
 - ▶ Deploying CUDA-accelerated applications
- ▶ Highly Evaluted

	benefit(5.0)	Total Evaluation(5.0)	Teaching Overall
Univ. Average	4.29	4.31	38.83
School. Average	4.29	4.31	38.83
This Course	4.50	4.44	40.61

- Comments from students:
 - Of High quality, learned a lot, very useful. Hope continuing.
 - Very practical.
- Some Other Universities Students went to USTC for it.



The Projects

Ray - Tracing

CUDA + Hadoop

Sparse Solver

3D Reconstruction

FDTD Simulation

Image&Video Processing

WPA2-PSK Attack

Speech Recognition

Page Rank, Search Engine

Machine Learning

Monte Carlo method

...

...



3D wave equation: >100x Speedup (dubious)


3D实现环境（服务器）

CPU	E5-2620
GPU	Tesla K20c
CUDA	5.0
3-D Grid	Support

三种方式的结果对比

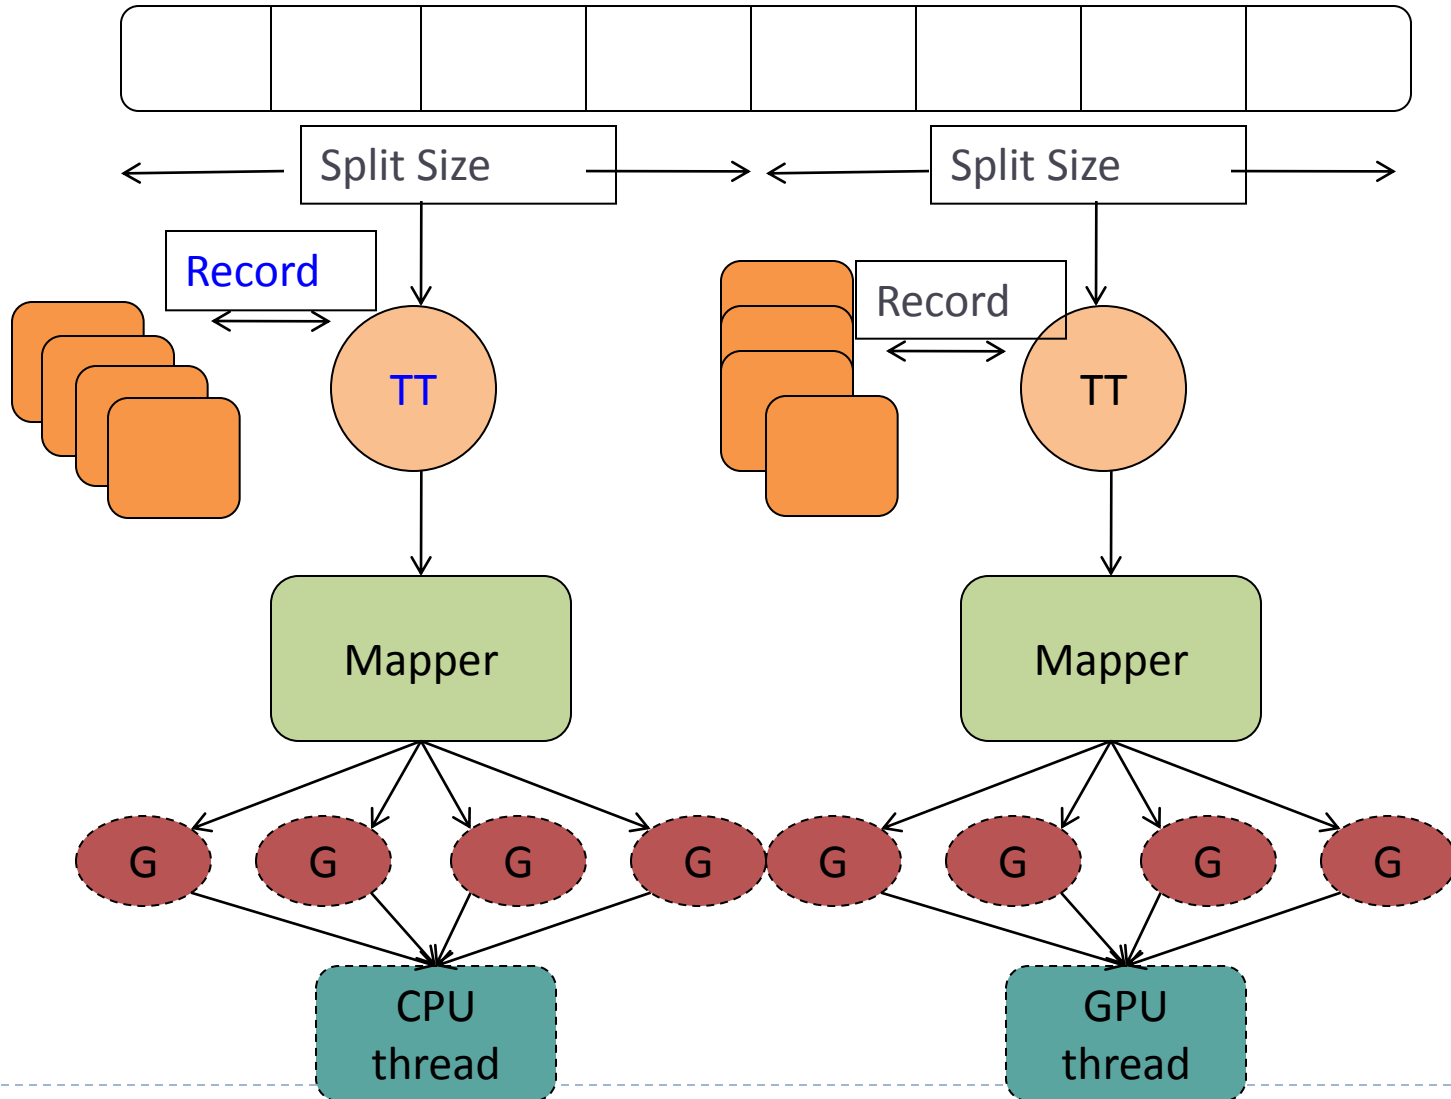
执行时间(ms)	
CPU	3170000
GPU3D	94.29084
GPU3D_SHARED	92.726463

z方向固定，k=128时，x, y, u的三维图

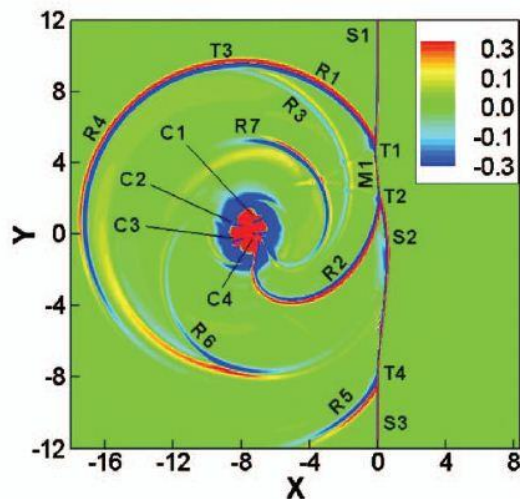



origin软件的分辨问题，在x,y方向上只取了64个点，进行作图，显得图形不圆滑，若是取所有点绘图，图形会非常圆滑


CUDA & Hadoop: Basic Function




FDTD Simulation: 10x Speedup



(h) $t = 12.0$



$t = 12.0$, single GPU

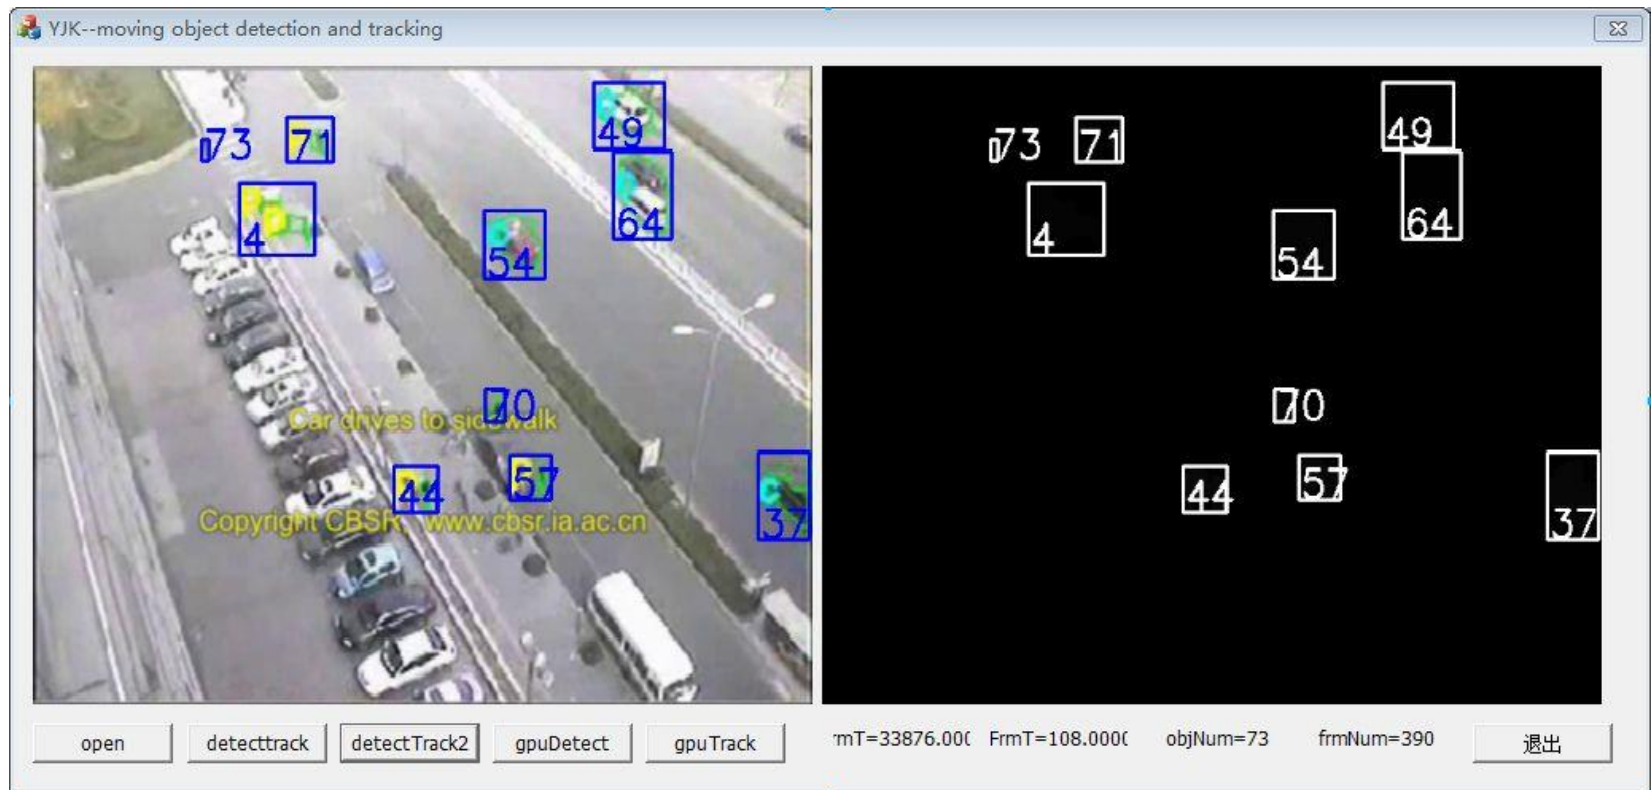


$t = 12.0$, dual GPUs


$t = 12.0$, reference

Video Multi-Target Tracking: Speedup 5x

Improvement needed



SIFT on GPU: 3x speedup



GPU-based Search Engine: 16x speedup



4 Students



30 Days



5modules



100,000 pages



1,219,508URLs



155 Kernel Codes



GPU 128 ms



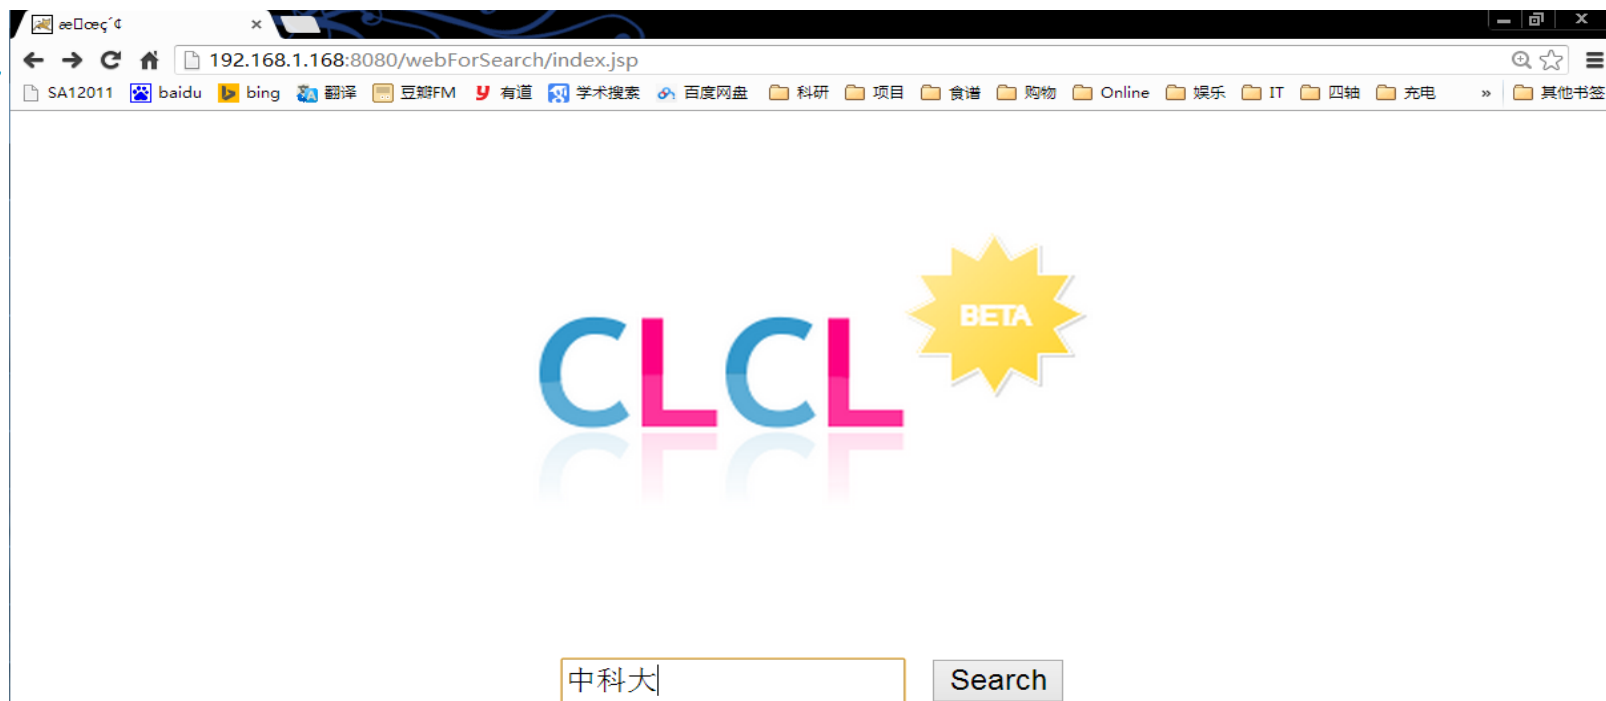
16x Vs. CPU



More ?



搜索首页



搜索结果

查询结果

192.168.1.168:8080/webForSearch/MyJsp.jsp

SA12011 baidu bing 翻译 豆瓣FM 有道 学术搜索 百度网盘 科研 项目 食谱 购物 Online 娱乐 IT 四轴 充电 其他书签

Search For 中科大 Find Total: 246total time for serching: 2ms

[中科大“学术猫”受捧 -新华财经-新华网](#)
PageRank=0.00003324
中科大“学术猫”受捧 -新华财经-新华网 中科大“学术猫”受捧 2013年11月15日
05:30:58 | 责任编辑: 李天真 | 来源: 光明网 11月12日, 中国科学院院士、**中科大**教授[more](#)

[如果在中科大遇见你 中国科学技术大学新闻网](#)
PageRank=0.00002995
如果在**中科大**遇见你 中国科学技术大学新闻网 如果在**中科大**遇见你 2013-11-01 逍遥派小掌门 如果在**中科大**遇见你, 在初冬的早晨, 要相约去吃一次食堂的早餐, 慢慢习惯阳光和你都在身旁, 想要[more](#)

[三千人腌250吨泡菜 -新华财经-新华网](#)
PageRank=0.00003324
中国新闻网 当地时间2013年11月13日, 韩国首尔, 韩国举行爱心泡菜分享活动, 超过三千名参加者现场组成爱心形状, 集体腌制250吨的泡菜。 1 1 支持键盘翻页←左右→ 分享到: 更多图片 **中科大**[more](#)


[稀世粉钻8300万美元 -新华财经-新华网](#)
PageRank=0.00003324
地时间11月13日, “粉红之星”粉钻在瑞士日内瓦拍卖, 这枚罕见的宝石最终以8300万美元成交, 写下史上宝石拍卖新高。图为“粉红之星”粉钻。(资料图) 1 1 支持键盘翻页←左右→ 分享到: 更多图片 **中科大**[more](#)

[\[科技日报\]接续断裂的创新链条— 中科大先研院为成果转化探路 中国科学技术大学新闻网](#)
PageRank=0.0000243
[科技日报]接续断裂的创新链条— **中科大**先研院为成果转化探路 中国科学技术大学新闻网 [科技日报]接续断裂的创新链条— **中科大**先研院为成果转化探路 2013-12-06 本报记者 吴长锋 11月11日[more](#)

[黄山雨后现“佛光” -新华财经-新华网](#)
PageRank=0.00002719
黄山风景区雨后放晴, 云海飘荡。黄山风景区始信峰出现今年入冬以来最为持久的“佛光”奇观, 令人叹为观止。中新社发方也广德 摄 图片来源: CNSPHOTO 1 1 支持键盘翻页←左右→ 分享到: 更多图片 **中科大**[more](#)

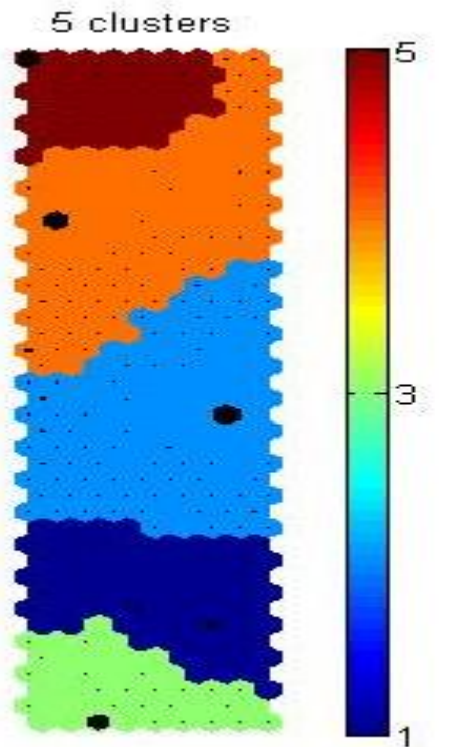
PCA on GPU: 5x Speedup

矩阵大小3072(样本数)*1024(特征数)




5x Speedup

GPU-SOM(Self Organizing Map): 50x

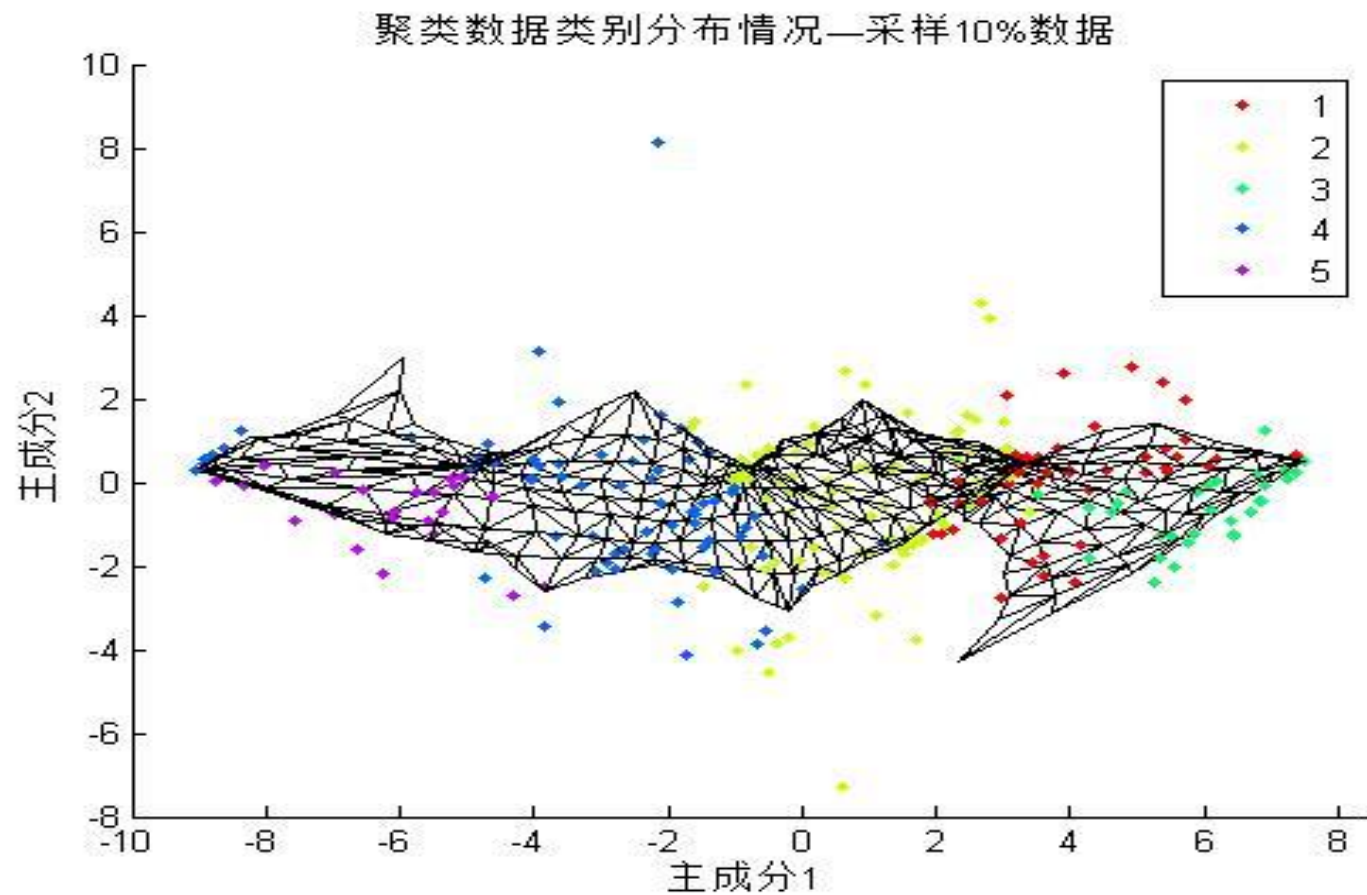


GPU



Matlab

Data Clustering



DeFrog: RealTime, 20x Speedup



Image Enhancement: 6x

1，硬件平台：

CPU-----Pentium 2.2GHz

GPU-----G210M,计算能力1.2


2，软件平台：

操作系统为 windows XP


运行环境为 VS2008， cuda

toolkit 5.5， opencv

3处理的图片的大小1200x800



	Data Transfer	Compute	Data Retrieve	Totally
CPU(ms)	14.211776	63.003777	11.579264	102.012772
GPU(ms)	1.878592	20.22376	2.439744	25.761248
GPU(ms)-Opt	1.878592	10.21354	2.439744	15.761248



- ▶ One of the world's Best SHA-1 on NVIDIA GPU

Performance Comparison

- **GTX 460**, 7 multiprocessors, 336 sp @ 1557MHz
 - ‘00000000’ – ‘00999999’, totally 10^6 messages
 - Our competitors are set into brute-force mode

	SHA-1	MD5
Our implementation (Hash/s)	400.3M	671.5M
oclHashcat-lite (Hash/s)	361.0M	1329.3M
IGHASHGPU (Hash/s)	371.5M	675.1M

课程资源

- ▶ NVIDIA CUDA Zone
 - ▶ <http://developer.nvidia.com/category/zone/cuda-zone>
- ▶ 网络教材和课程
 - ▶ <http://www.bb.ustc.edu.cn/>
- ▶ 相关素材
 - ▶ 将通过网络和电子邮件发布
- ▶ Others
 - ▶ 请各位多利用网络搜索引擎



预修课程和教材

▶ 基础知识

- ▶ 1) 计算机体系结构基础 2) C语言程序设计
- ▶ 3) 计算机算法基础 4) 线性代数

▶ 课程内容参考

- ▶ 1. CUDA C Programming Guide, NVIDIA Corp.
- ▶ 2. CUDA Best Practice Guide, NVIDIA Corp.
- ▶ 3. Programming Massively Parallel Processors, 2010, David Kirk and Wen-mei Hwu

▶ 声明：参考和引用了如下的工作

- ▶ 1. Patrick Cozzi, CIS 565, University of Pennsylvania
- ▶ 2.
- ▶ 3.
- ▶ 4. Udacity CS 344 Intro to Parallel Computing



Know the backgrounds and computer level of the students



- Many Students don't have solid programming backgrounds
- Adding Some contents about C programming will be helpful
- **Add Some Basic Computer Architecture Contents**
- **Add Some Basic numerical methods**
- **Add some Parallel Computing contents**



Prerequisite

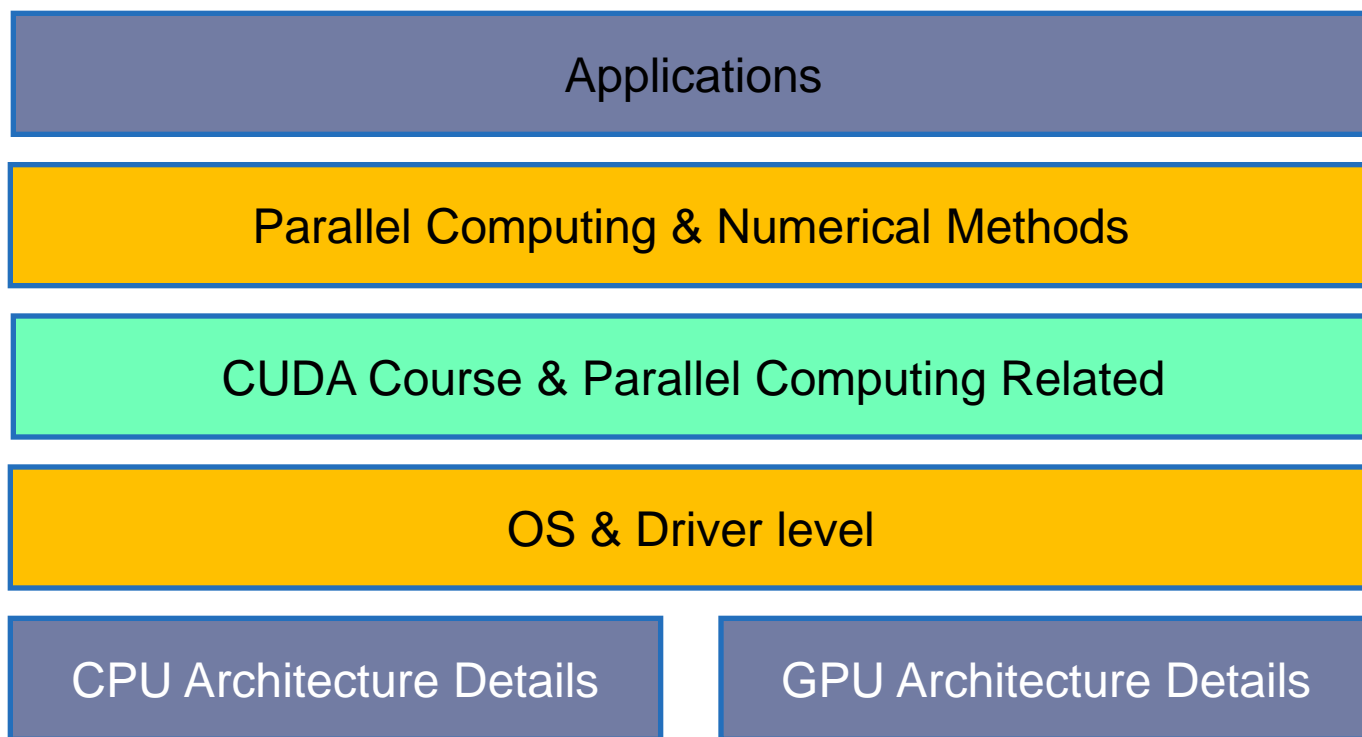
- C Programming Language
- Basic Computer Architecture
- Basic Algebra and Numerical Methods
- Domain Knowledge
- Linux OS



Prerequisite

- Many students have high scores but not really know how and why
- Don't take Prerequisite seriously when they claim they are ok. They are really not!
- However, give them basic trainings and review on the prerequisites

课程定位



课程目标

- ▶ 计算机，电子，自动化，生医等相关专业的硕士研究生或高年级本科生
- ▶ 了解和掌握GPU并行计算系统的分析，设计，开发，调试和优化方法
- ▶ 系统分析能力，编程能力，开发经验





目标说明

- Not theoretical Algebra or Numerical methods
 - ⑩ But needs to add some basic algebra
- Not Computer Architecture Course
 - ⑩ But needs to add some basic contents
- Not Parallel Computing Course
 - ⑩ But needs to add some contents

课程模式

- ▶ 强调面向应用和实践
- ▶ 最新的业界动态和技术
- ▶ 课堂讲解+上机实验
- ▶ 评分方式
 - ▶ 1) 平时作业和实验报告 40%
 - ▶ 2) 课程考试 25%
 - ▶ 3) 大作业 35%
- ▶ 学生小组制




作业提交

- ▶ 每周六，周日 讲完课程布置作业
- ▶ 在下一个周四24:00之前提交
- ▶ 提交方式
 - ▶ 网络系统：以服务器时间为准
 - ▶ 电子邮件：以我或助教的电脑时间为准
 - ▶ 邮寄方式：以当地邮局邮戳为准
 - ▶ 当面提交：以我或助教的手机手表时间为准
- ▶ 晚交一天扣本次作业10%的分，直至扣完
- ▶ 严禁抄袭，凡抄袭和被抄袭者都是0分



开始课程之前的问题…

- ▶ 1) 为什么我们要使用计算机?
 - ▶ 为了更好地解决计算问题
- ▶ 2) 你需要什么样的计算机? 畅想…
 - ▶ 速度更快
 - ▶ 内存更大
 - ▶ 无穷多的外存
 - ▶ 智能化的接口



可惜世间总是太多无奈...

- ▶ 现在的CPU系统已经遇到各种瓶颈
- ▶ 只能向多核及并行系统发展
- ▶ 顺势而生的 *GPU* – Graphics *P*rocessing *U*nit



内容

- ▶ 1) CPU体系架构概述
 - ▶ 2) 并行程序设计概述
 - ▶ 3) CUDA开发环境搭建和工具配置
 - ▶ 4) GPU体系架构概述
 - ▶ 5) GPU编程模型
 - ▶ 6) CUDA编程(1)
 - ▶ 7) CUDA编程 (2)
 - ▶ 8) CUDA编程 (3)
 - ▶ 9) CUDA程序分析和调试工具
 - ▶ 10) CUDA程序基本优化
 - ▶ 11) CUDA程序深入优化
 - ▶ 12) 最新NVIDIA GPU 和 CUDA特性
-





课程内容

- a) what to cover & what not to?
- b) GPU Architecture & CPU Architecture ?
- c) CUDA Programming
- d) CUDA Optimizations
- e) advance Topics : MPI + CUDA, Multi-GPU, Clusters, Numerical Algorithms with GPU, OpenCL, OpenGL+CUDA, Cloud Computing



What to cover

- Basic C Programming & Linux OS
- Basic Computer Architecture
- Basic Parallel Computing
- Basic CPU Architecture
- Basic GPU Architecture
- CUDA Programming Language
- CUDA Optimization
- MPI+OpenMP
- Compiling & Debugging

What to cover

- Tools + Third Party Software
 - ⑩ Compiling
 - ⑩ Debugging Tools
 - ⑩ Nsight, CUDA GDB
 - ⑩ Visual Profiler
- Needed Linux OS



What **NOT** to cover?

- GPU architecture Details
- Pipeline
- Programmable Vertex and Pixel Processors
- Texture Details
- Low-level drivers

GPU Architecture & CPU Architecture ?



- Basic Information will be enough
- ALU
- Register
- Data Path
- Cache
- Instruction Execution
- SIMD+SIMT
- SM,SMX Info



CUDA Programming

- C Programming Improvements
- Basic Programming Guide
- Cover as much as possible
- No low-level programming
 - ⑩ Such as PTX, Assemble...
- 10 hours of class



CUDA Programming cont' d

- More Practice based
- Every class is followed by lab work
 - ⑩ Lab work including whole development process
- Every class is followed by a small programming project
 - ⑩ Vector Add
 - ⑩ Parallel reduction
 - ⑩ Matrix Multiplication
 - ⑩ And so on...
- Debugging & Profiling are to be added
- Tools
- More TA, more basic programming guidance.



CUDA Optimizations

- Know how and why
- Focusing on bottleneck analysis
- More bandwidth-limited applications
- Coalesced Memory Access
- Using Visual Profiler to find problems
- Emphasize Algorithm Design
- Not focusing on Program Optimization



Advance Topics

- MPI + CUDA,
- MultiGPU
- Remote GPU Clusters
 - ⑩ To promote after-class study
- Numerical Algorithms with GPU
 - ⑩ Numerical Errors
- OpenCL
- OpenGL+CUDA
- Cloud Computing & Grid




Research Fields Related Contents

- Asking Students to participate in teaching
- Students work within 3~5 Groups
- Every group raise a specific topic
- Students give talks
- More research activities involvements

课程内容的关系

- ▶ 应用是最关键的
 - ▶ Application Driven!!!
- ▶ 并行处理: 了解架构
- ▶ 编程
- ▶ 优化




课程内容

► WebGL 展示



WebGL Skin

http://alteredqualia.com/three/examples/webgl_materials_skin.html



WebGL Water

<http://madebyevan.com/webgl-water/>

课程内容

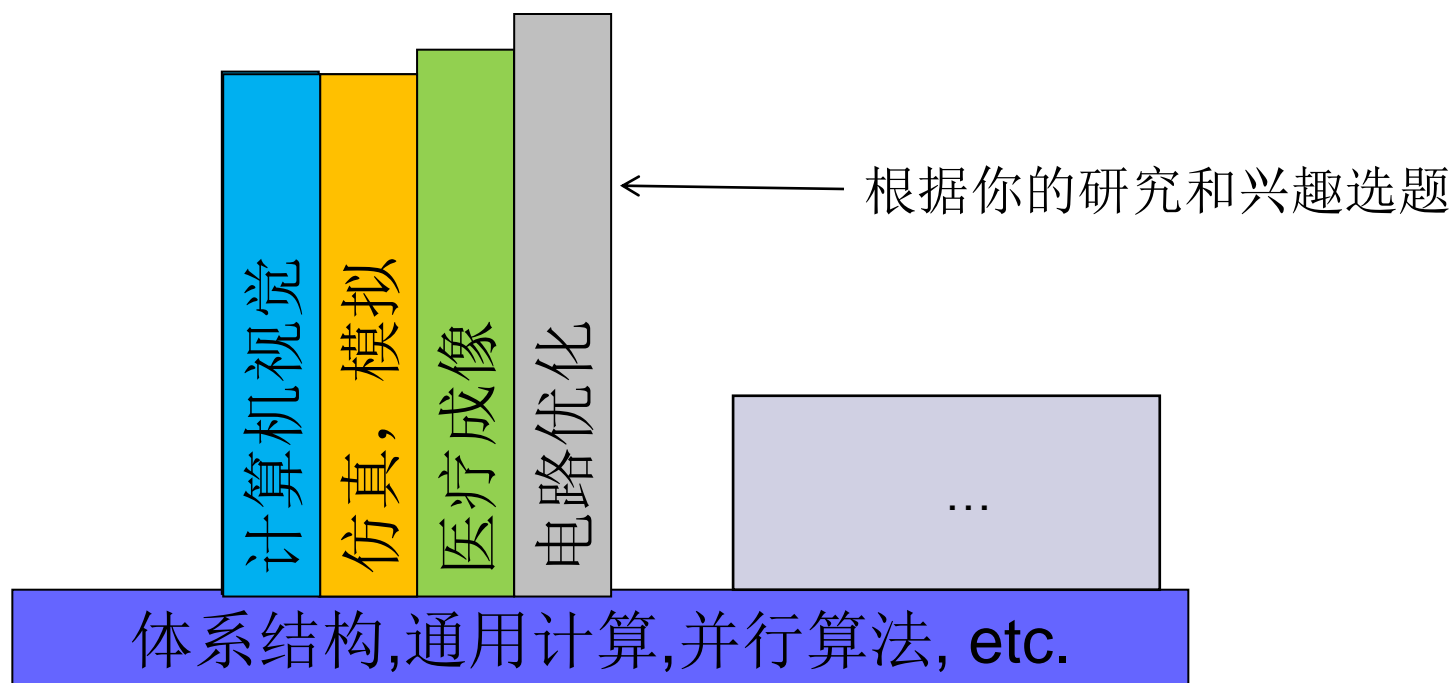
► GPU Compute + Real-Time Rendering



http://www.nvidia.com/object/GTX_400_games_demos.html

课程内容

► 大作业:



硬件需求

- ▶ 上机实验需要具有 *NVIDIA Fermi* 以上的GPU
- ▶ 作业将同样需要,可以到实验室使用
- ▶ 更新驱动
 - ▶ <http://www.nvidia.com/Download/index.aspx>
- ▶ 如何看GPU型号?



软件和后续学习

- ▶ 将学习使用Windows和 Linux 系统
- ▶ 喜欢Linux系统
 - ▶ 相信科大同学的能力
- ▶ 联系助教获得其他实验室资源
- ▶ 联系教师获得其他资源
- ▶ 实验地点：电四楼 209 实验室

